

Integrating OpenTelemetry & Security in eShop

Daniel Madureira

N^o.Mec: 107603

Professor Cláudio Teixeira

Software Architectures

DETI

Universidade de Aveiro

March 2025

Contents

1	Introduction	2
2	Implementation	2
2.1	Traces with Jaeger	3
2.1.1	Configuring Jaeger	3
2.1.2	Configure traces	3
2.2	Metrics with Prometheus	7
2.2.1	Configure Prometheus	7
2.2.2	Configure Custom Metrics	8
2.3	Masking sensitive data	11
2.4	Visualization with Grafana	14
2.4.1	Dashboards	15
3	Conclusion	17

1 Introduction

The purpose of this assignment was to implement observability using traces and metrics for a single feature in an eShop application. The objectives can be summarized as follows:

- **Implement OpenTelemetry tracing** for a specific feature or use case (end-to-end).
- **Mask or exclude sensitive data** (e.g., email, payment details) from telemetry and logs.
- Set up a basic **Grafana dashboard** to visualize **traces** and **metrics**.

For this, we used a full-stack eShop application written in .NET, available [here](#). Since some modifications were necessary to run the code locally on Fedora 41, the instructions for running the application, along with the code implementing my solution, are available [here](#).

This report will explore the developed solution in more detail.

2 Implementation

As indicated in the assignment, I chose the *place an order from basket* use case to implement tracing and metrics. To achieve this, I used **Jaeger** for tracing, **Prometheus** for metrics, and the **OpenTelemetry Collector** to gather metrics via the OpenTelemetry protocol.

With the use case defined and the tech stack selected, the next step was to identify which microservices were involved in this workflow. To do this, I utilized the built-in traces dashboard:

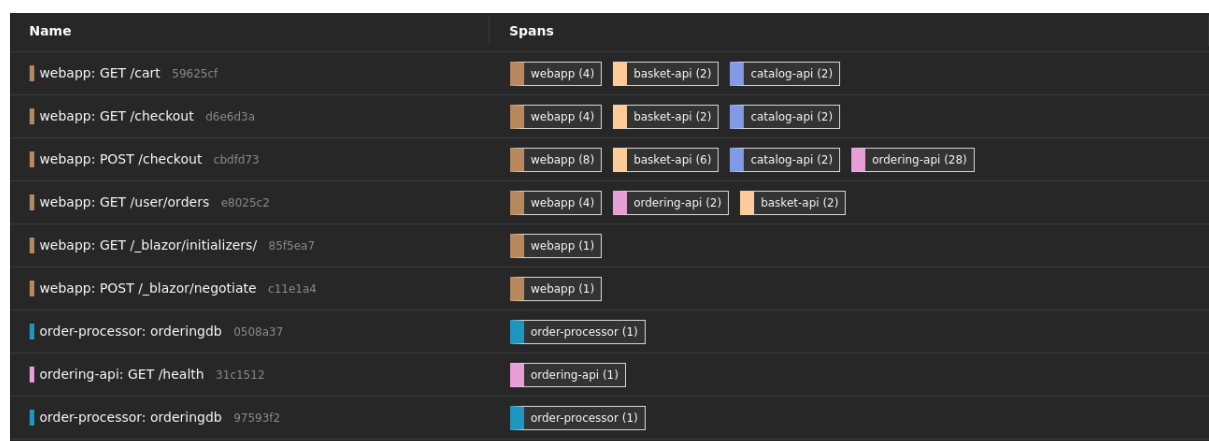


Figure 1: Microservices involved in the use case

From this, I determined that traces and metrics needed to be integrated into the following microservices: **Basket.API**, **Ordering.API**, **WebApp**, and **Order-Processor**.

2.1 Traces with Jaeger

2.1.1 Configuring Jaeger

As mentioned before, I chose **Jaeger** to receive and view traces from the microservices. To achieve this, I used the **Jaeger Docker image** and included it in a `docker-compose.yml` file as follows:

```
services:
  jaeger:
    image: jaegertracing/all-in-one:latest
    container_name: jaeger
    ports:
      - "16686:16686" # Jaeger UI
      - "4317:4317"   # OTLP gRPC
      - "4318:4318"   # OTLP HTTP
```

With the **image** defined and the **container** running under the name *jaeger*, I can now access the **Jaeger UI** on **port 16686**. Additionally, I have defined **port 4317** for **OTLP** communication over **gRPC** and **port 4318** for **OTLP** communication over **HTTP**.

Since I didn't specify a custom network, all services inside this **Docker Compose** file are built within the **default network**, allowing them to communicate with each other.

By running the **Docker Compose** file, we can finally view the **Jaeger interface**.




Figure 2: Jaeger UI

2.1.2 Configure traces

Now that we have **Jaeger** running and working, we need to configure the microservices to send traces to it. To do this, I used **OpenTelemetry** to export traces over **gRPC** to **Jaeger** by modifying the `Program.cs` file in each microservice.

For example, in the **Basket.API** microservice, I added the following code:

```

const string serviceName = "eShop.Basket.API";

// Configure OpenTelemetry
builder.Services.AddOpenTelemetry()
    .WithTracing(
        (tracing) =>
        {
            tracing.AddSource(serviceName);
            tracing.AddProcessor(new MaskProcessor());
            tracing.AddAspNetCoreInstrumentation();
            tracing.AddHttpClientInstrumentation();
            tracing.AddGrpcClientInstrumentation();
            tracing.AddOtlpExporter(
                options =>
                {
                    options.Endpoint = new Uri("http://localhost:4317");
                    options.Protocol = OtlpExportProtocol.Grpc;
                }
            );
        }
    );

```

Looking deeper into the code, we can see where we:

- **Configure OpenTelemetry Tracing:** The `.WithTracing` method is used to set up tracing, enabling the collection and recording of application execution details, such as requests, responses, and events.
- **Add a Telemetry Source:** The `tracing.AddSource(serviceName)` method identifies the service that is generating telemetry data, helping to distinguish between different data sources.
- **Mask Sensitive Data:** By using `tracing.AddProcessor(new MaskProcessor())`, a custom processor is added to filter or mask sensitive information before it is exported, ensuring data privacy.
- **Enable Instrumentation:**
 - `tracing.AddAspNetCoreInstrumentation()` captures details of incoming HTTP requests in ASP.NET Core applications.
 - `tracing.AddHttpClientInstrumentation()` tracks and records outgoing HTTP requests.
 - `tracing.AddGrpcClientInstrumentation()` captures traces of gRPC client calls.
- **Configure OTLP Exporter:** The `tracing.AddOtlpExporter` method configures the exporter to send telemetry data to an OTLP collector over gRPC at `http://localhost:4317`, ensuring that the data is transmitted accurately for storage and analysis.

This ensures that the **Basket.API** service collects and forwards data to **Jaeger**. After this, I added custom traces to the services. I will showcase the work done on the **Basket.API**, but the other microservices were configured using the same principles.

In this case, I created a custom **ActivitySource** with the same **serviceName** as the one configured inside the **Program.cs** file:

```
private static readonly ActivitySource ActivitySource = new
    ActivitySource("eShop.Basket.API");
```

This **ActivitySource** is used to create and manage tracing activities within the **BasketService** class. Inside each function (**GetBasket**, **UpdateBasket**, and **DeleteBasket**), a new activity is started using:

```
using var activity = ActivitySource.StartActivity("GetBasket",
    ActivityKind.Server);
```

We then collect **custom traces** by setting **tags** as follows:

```
activity?.SetTag("user.id", context.GetUserIdentity());
activity?.SetTag("request.method", context.Method);

var userId = context.GetUserIdentity();
if (string.IsNullOrEmpty(userId))
{
    activity?.SetStatus(ActivityStatusCode.Error, "User ID is null or
        empty");
    return new();
}

if (logger.IsEnabled(LogLevel.Debug))
{
    logger.LogDebug("Begin GetBasketById call from method {Method} for
        basket id {Id}", context.Method, userId);
}

var data = await repository.GetBasketAsync(userId);

if (data is not null)
{
    activity?.SetStatus(ActivityStatusCode.Ok, "Basket retrieved
        successfully");
    return MapToCustomerBasketResponse(data);
}

activity?.SetStatus(ActivityStatusCode.Error, "Basket not found");
return new();
```

In this case, we:

- Add Tags:

- `activity?.SetTag("user.id", context.GetUserIdentity())` attaches the user identity to the trace for better context.
 - `activity?.SetTag("request.method", context.Method)` adds the gRPC method name, providing additional context for the trace.
- **Set Status:**
 - If the user ID is **null or empty**, the trace status is marked as **Error**.
 - If the basket is **successfully retrieved**, the status is set to **OK**.
 - If the basket is **not found**, the trace status is marked as **Error**.

These traces can then be visualized in **Jaeger UI** under **basket-api**.

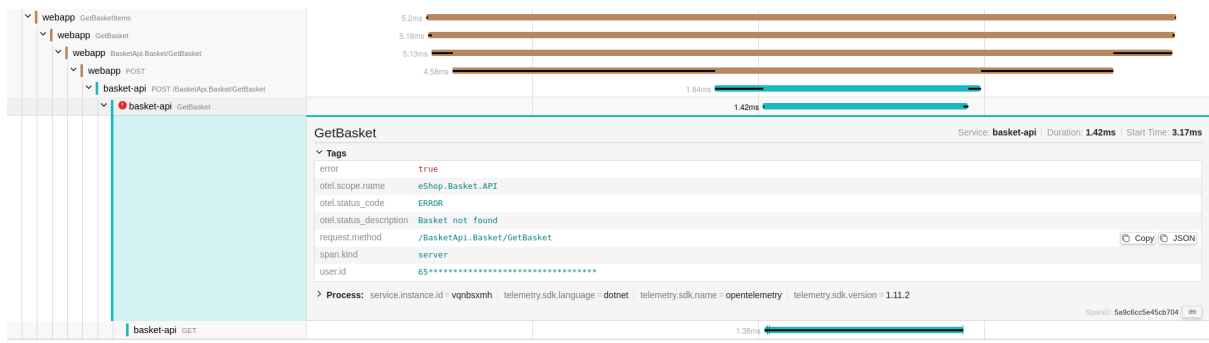


Figure 3: GetBasket traces

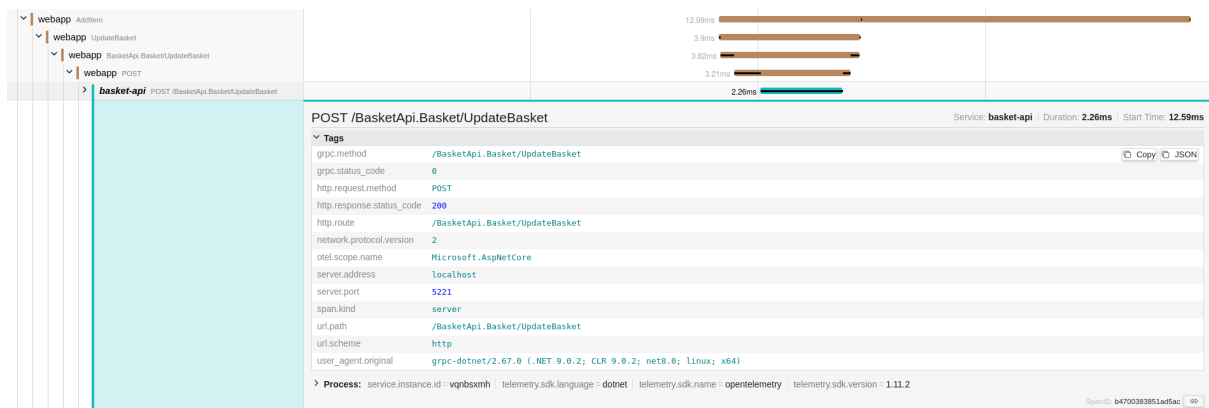


Figure 4: UpdateBasket traces



Figure 5: DeleteBasket traces

2.2 Metrics with Prometheus

2.2.1 Configure Prometheus

In order to collect metrics, I chose to use **Prometheus** and **OpenTelemetry Collector** (**otel-collector**) to collect and store the metrics defined in the application. For this purpose, I created two configuration files, **prometheus.yml** and **otel-collector/config.yml**, that hold the configuration for these two services.

The **otel-collector** is responsible for collecting metrics from the different services and exporting them to **Prometheus**, acting as a middleman. Without the **otel-collector**, Prometheus would need to pull the metrics from each service individually. However, by using the **otel-collector**, we can collect, process, and send metrics to Prometheus in a more structured and efficient way. In this case, the **otel-collector** receives data over **gRPC** on port **4316** or **HTTP** on port **4315**, and exports them to the Prometheus endpoint on port **8889**:

```
#otel-collector/config.yml
receivers:
  otlp:
    protocols:
      grpc:
        endpoint: 0.0.0.0:4316
      http:
        endpoint: 0.0.0.0:4315

exporters:
  prometheus:
    endpoint: 0.0.0.0:8889

service:
  pipelines:
    metrics:
      receivers: [otlp]
      exporters: [prometheus]
```

On the Prometheus side, it is configured to collect data every **15 seconds** from the defined endpoint. In this case, Prometheus collects data from the **otel-collector** on port **8889**. The configuration for this is defined in the **prometheus.yml** file:


```
#prometheus.yml
global:
  scrape_interval: 15s

scrape_configs:
  - job_name: 'otel-collector'
    static_configs:
      - targets: ['otel-collector:8889']
```

Similar to the Jaeger setup, I used the respective **Docker images** to launch these services on the same network as Jaeger by adding them to the `docker-compose.yml` file. Here's the configuration:

```
# ... jaeger service
otel-collector:
  container_name: otel-collector
  image: otel/opentelemetry-collector-contrib
  volumes:
    - ./observability/otel-collector/config.yaml:
        ↪ /etc/otelcol-contrib/config.yaml
  ports:
    - 1888:1888 # pprof extension
    - 8888:8888 # Prometheus metrics exposed by the Collector
    - 8889:8889 # Prometheus exporter metrics
    - 13133:13133 # health_check extension
    - 4316:4316 # OTLP gRPC receiver
    - 4315:4315 # OTLP http receiver
    - 55679:55679 # zpages extension
  depends_on:
    - prometheus

prometheus:
  container_name: prometheus
  image: prom/prometheus:latest
  volumes:
    - ./observability/prometheus/prometheus.yml:
        ↪ /etc/prometheus/prometheus.yml
  ports:
    - "9090:9090"
```

After relaunching the Docker Compose file, we can now see two additional services running: the **otel-collector** service is running on port **4316** to receive metrics over **gRPC**, and on port **8889** to export them to Prometheus. Additionally, the **Prometheus UI** is now accessible on port **9090**, where we can view and analyze the collected metrics.

2.2.2 Configure Custom Metrics

Similar to adding tracing to the microservices, I had to modify the `Program.cs` file in each microservice to enable the collection of metrics. Once again, I will use the `Basket.API`

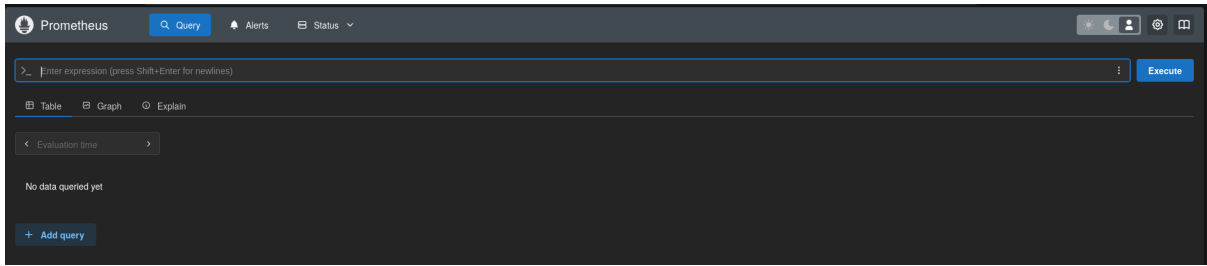


Figure 6: Prometheus UI

microservice as an example, but the other microservices in the use case had similar changes made to them.

To collect the metrics, I added the following code to `Program.cs`:

```
// builder.Services.AddOpenTelemetry()
// .WithTracing(
//     ...tracing config
// )
.WithMetrics(
    metrics =>
    {
        metrics.AddHttpClientInstrumentation();
        metrics.AddAspNetCoreInstrumentation();
        metrics.AddMeter(serviceName, serviceVersion);
        metrics.AddOtlpExporter(
            options =>
            {
                options.Endpoint = new Uri("http://localhost:4316");
            }
        );
    }
)
```

Looking at the changes, we can identify the following:

- **Metrics Collection Setup:** Initializes OpenTelemetry to collect performance metrics such as latency, errors, and throughput.
- **Outgoing HTTP Request Metrics:** Tracks external HTTP requests, capturing metrics like request count, response time, and errors.
- **ASP.NET Core HTTP Request Metrics:** Captures incoming HTTP request data, including count, duration, and response status.
- **Meter Setup:** Records metrics tied to a specific service and version, providing better tracking.
- **OTLP Exporter Configuration:** Sends collected metrics to the `otel-collector` via OTLP for analysis and processing.

On the other side, I added custom metrics to the `BasketService.cs` using counters and histograms. The following code defines these metrics:

```
private static readonly Meter Meter = new Meter("eShop.Basket.API",
    ↪ "1.0.0");
private static readonly Counter<int> GetBasketCounter =
    ↪ Meter.CreateCounter<int>("get_basket_requests");
private static readonly Counter<int> UpdateBasketCounter =
    ↪ Meter.CreateCounter<int>("update_basket_requests");
private static readonly Counter<int> DeleteBasketCounter =
    ↪ Meter.CreateCounter<int>("delete_basket_requests");
private static readonly Histogram<double> RequestDurationHistogram =
    ↪ Meter.CreateHistogram<double>("request_duration", "ms", "Duration of
    ↪ requests in milliseconds");
```

Here, I defined a meter group under the service name `eShop.Basket.API` and created counters to track the number of times requests are made to get, update, or delete the basket. I also created a histogram to measure the duration of requests. These counters and histograms are updated in their respective functions. For example, in the `GetBasket` function:

```
public override async Task<CustomerBasketResponse>
    ↪ GetBasket(GetBasketRequest request, ServerCallContext context)
{
    GetBasketCounter.Add(1);

    // ... define traces

    var stopwatch = Stopwatch.StartNew();

    // ... get basket logic

    stopwatch.Stop();
    RequestDurationHistogram.Record(stopwatch.ElapsedMilliseconds, new
    ↪ KeyValuePair<string, object>("method", "GetBasket"));

    // ... rest of the function
}
```

With everything configured, when we access the **Prometheus UI** and explore the available metrics, we can see both our custom metrics and some default metrics from `ASP.NET`.

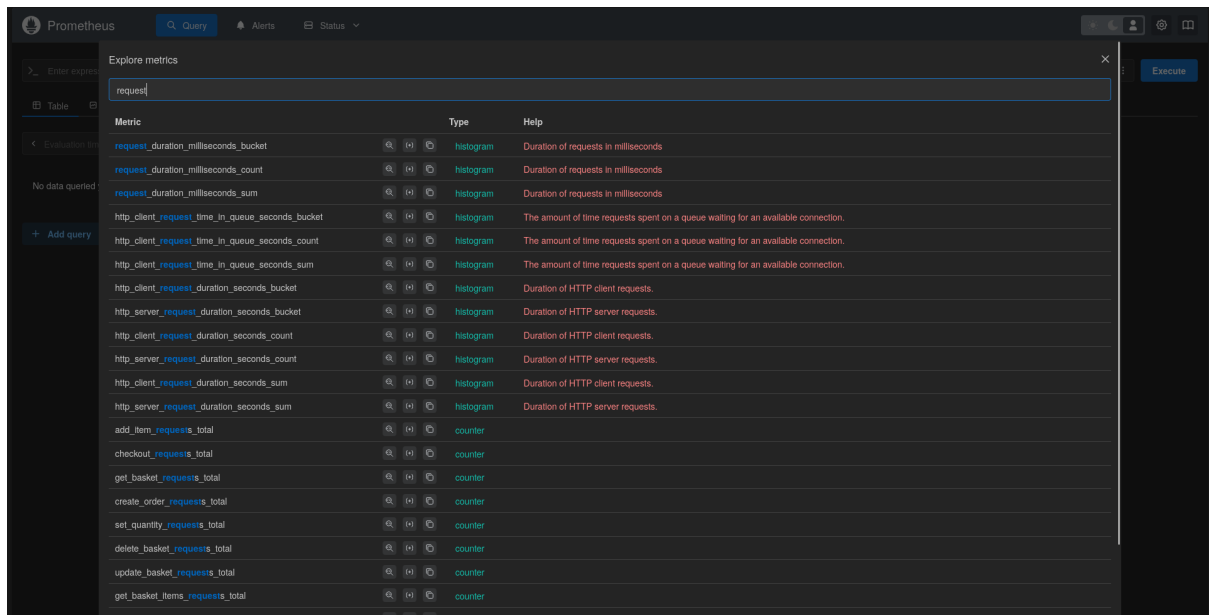


Figure 7: Metrics from Prometheus

2.3 Masking sensitive data

As my use case involves sensitive information, such as payment details like card numbers and security codes, there was a need to either avoid tracing this kind of data or at least mask it in the traces. To achieve this, when exporting metrics to Jaeger, I added an `ActivityProcessor` to mask sensitive data.

Here is the code for the custom `MaskProcessor`:

```
public class MaskProcessor : BaseProcessor<Activity>
{
    public override void OnEnd(Activity activity)
    {
        if (activity.Tags.Any(tag => tag.Key == "user.id"))
        {
            var userId = activity.Tags.FirstOrDefault(tag => tag.Key ==
                "user.id").Value;
            var maskedUserId = MaskString(userId);
            activity.SetTag("user.id", maskedUserId);
        }
        if (activity.Tags.Any(tag => tag.Key == "request.command.buyer"))
        {
            var buyer = activity.Tags.FirstOrDefault(tag => tag.Key ==
                "request.command.buyer").Value;
            var maskedBuyer = MaskString(buyer);
            activity.SetTag("request.command.buyer", maskedBuyer);
        }
        if (activity.Tags.Any(tag => tag.Key ==
            "request.command.buyerId"))
        {

```

```

    var buyerId = activity.Tags.FirstOrDefault(tag => tag.Key ==
        ↪ "request.command.buyerId").Value;
    var maskedBuyerId = MaskString(buyerId);
    activity.SetTag("request.command.buyerId", maskedBuyerId);
}
if (activity.Tags.Any(tag => tag.Key ==
    ↪ "request.command.cardNumber"))
{
    var cardNumber = activity.Tags.FirstOrDefault(tag => tag.Key
        ↪ == "request.command.cardNumber").Value;
    var maskedCardNumber = MaskString(cardNumber);
    activity.SetTag("request.command.cardNumber",
        ↪ maskedCardNumber);
}
if (activity.Tags.Any(tag => tag.Key ==
    ↪ "request.command.cardHolderName"))
{
    var cardHolderName = activity.Tags.FirstOrDefault(tag =>
        ↪ tag.Key == "request.command.cardHolderName").Value;
    var maskedCardHolderName = MaskString(cardHolderName);
    activity.SetTag("request.command.cardHolderName",
        ↪ maskedCardHolderName);
}
if (activity.Tags.Any(tag => tag.Key ==
    ↪ "request.command.cardExpiration"))
{
    var cardExpiration = activity.Tags.FirstOrDefault(tag =>
        ↪ tag.Key == "request.command.cardExpiration").Value;
    var maskedCardExpiration = MaskString(cardExpiration);
    activity.SetTag("request.command.cardExpiration",
        ↪ maskedCardExpiration);
}
if (activity.Tags.Any(tag => tag.Key ==
    ↪ "request.command.cardSecurityNumber"))
{
    var cardSecurityNumber = activity.Tags.FirstOrDefault(tag =>
        ↪ tag.Key == "request.command.cardSecurityNumber").Value;
    var maskedCardSecurityNumber = MaskCVV(cardSecurityNumber);
    activity.SetTag("request.command.cardSecurityNumber",
        ↪ maskedCardSecurityNumber);
}
if (activity.Tags.Any(tag => tag.Key ==
    ↪ "request.command.Street"))
{
    var street = activity.Tags.FirstOrDefault(tag => tag.Key ==
        ↪ "request.command.Street").Value;
    var maskedStreet = MaskString(street);
    activity.SetTag("request.command.Street", maskedStreet);
}

```

```

    }
    if (activity.Tags.Any(tag => tag.Key ==
        ↪ "request.command.ZipCode"))
    {
        var zipCode = activity.Tags.FirstOrDefault(tag => tag.Key ==
            ↪ "request.command.ZipCode").Value;
        var maskedZipCode = MaskString(zipCode);
        activity.SetTag("request.command.ZipCode", maskedZipCode);
    }
}

private static string MaskString(string str)
{
    return str.Substring(0, 2) + new string('*', str.Length - 2);
}

private static string MaskCVV(string str)
{
    return new string('*', str.Length);
}
}

```

This class processes telemetry data, specifically activities, and masks sensitive information before it is exported. The `OnEnd` method is overridden to apply the masking operations when an activity ends. Inside the `OnEnd` method, we define the tags containing sensitive information and apply a mask accordingly.

The `MaskProcessor` is then added to the tracing configuration in the `Program.cs` file of each microservice, as shown below:

```

// config openTelemetry
builder.Services.AddOpenTelemetry()
    .WithTracing(
        (tracing) =>
        {
            \N ... rest of implementation

            tracing.AddProcessor(new MaskProcessor());

            \N ... rest of implementation
        }
    );

```

As a result, we can confirm that the sensitive fields are masked in the traces, ensuring that no sensitive information is exposed:

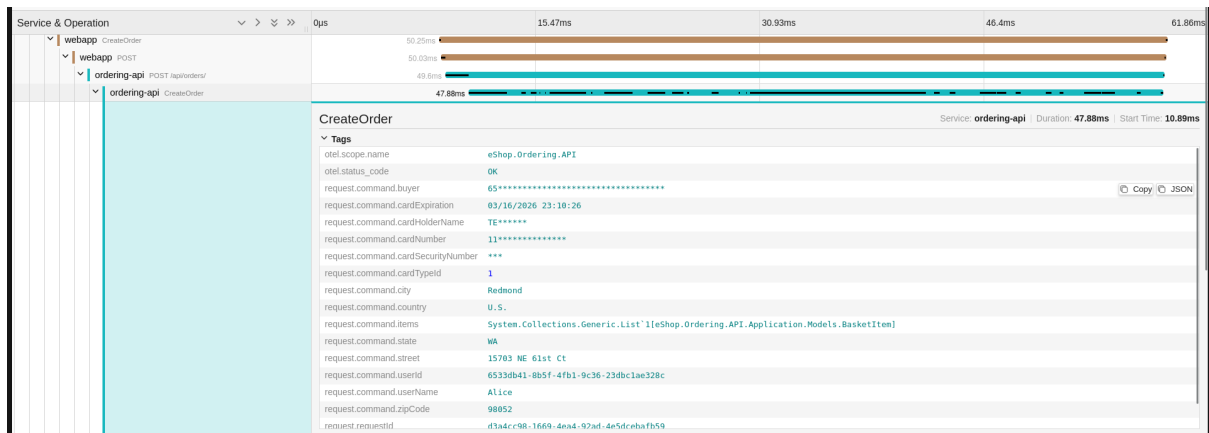


Figure 8: Masked fields in the request

2.4 Visualization with Grafana

After collecting traces and metrics, we need to create a dashboard to visualize the data over time as users interact with the application. To achieve this, I used **Grafana**. Once again, I leveraged Docker images and added them to the Docker Compose file:

```
services:
  # ... rest of services
  grafana:
    container_name: grafana
    image: grafana/grafana:latest
    ports:
      - "3000:3000"
    environment:
      - GF_SECURITY_ADMIN_PASSWORD=admin
    volumes:
      - ./observability/grafana/provisioning:
          ↪ /var/lib/grafana/provisioning
      - ./observability/grafana/dashboards:
          ↪ /var/lib/grafana/dashboards
```

When we run the Docker Compose file, **Grafana** will be running on **port 3000**. It connects to **Prometheus**, **Jaeger**, and **otel-collector** to display the metrics and traces collected by these services in custom dashboards. By default, **Grafana** comes with an admin login, with both the **username** and **password** set to 'admin'. It also maps volumes, though it doesn't provide default dashboards. However, Grafana can import **JSON files** with pre-configured dashboards.

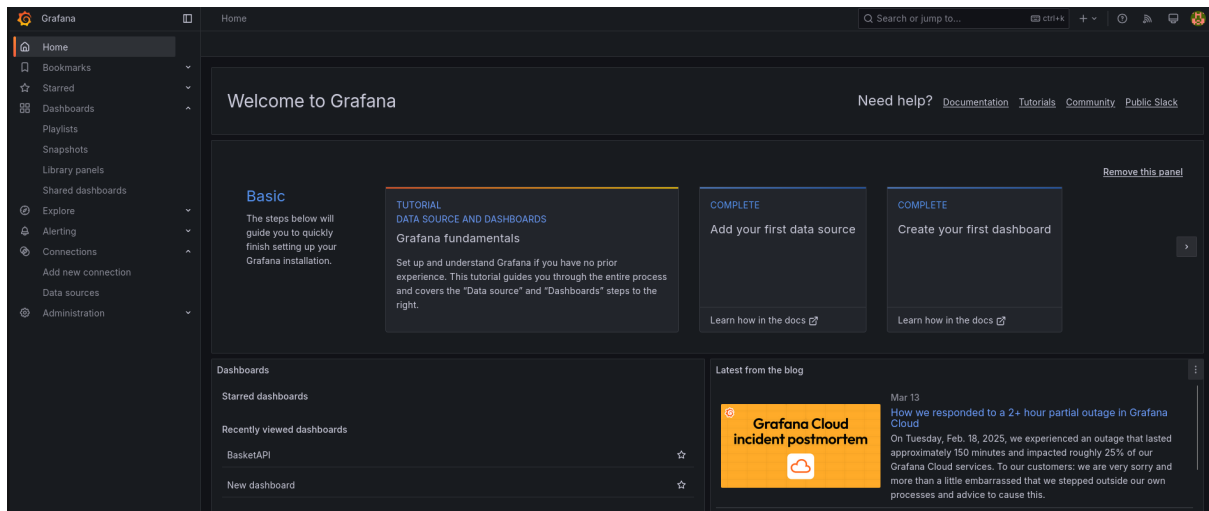


Figure 9: Grafana UI

To create a new dashboard, we first need to configure the **data sources** by adding the **Prometheus** and **Jaeger** URLs:

- `http://prometheus:9090`
- `http://jaeger:16686`

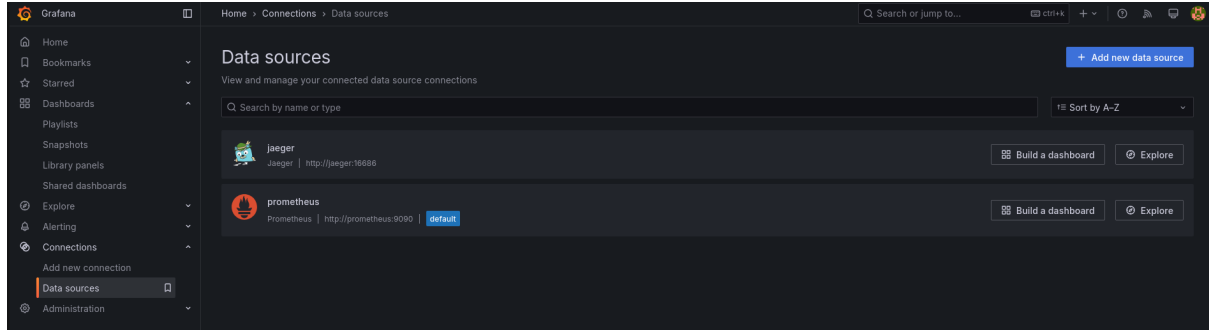


Figure 10: Data sources in Grafana

2.4.1 Dashboards

Having the data flowing from the sources, I built a dashboard to display some metrics and traces that the services collect over time for two microservices: **Basket.API** and **Ordering.API**:

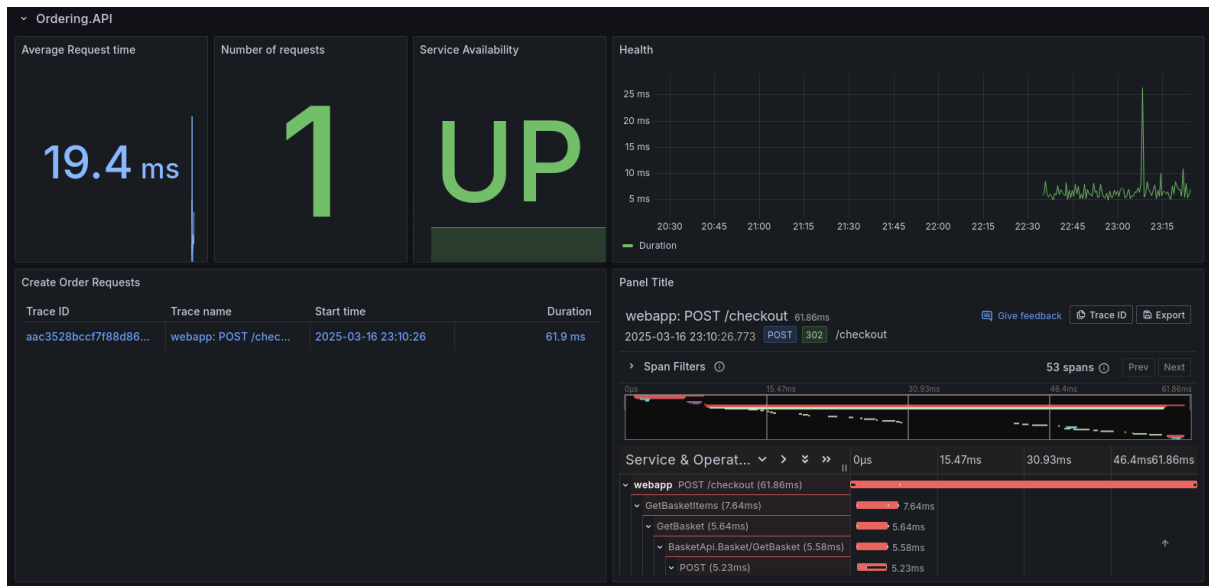


Figure 11: Ordering dashboard

Here, we display the following:

- **Average request time:** The average time for each request to the Ordering API.
- **Number of requests:** This metric counts the number of requests made to the Ordering API, collected by **Jaeger**.
- **Service availability:** This data is gathered from the health checks in **ASP.NET**.
- **Health:** A Time Series Graph that displays the API response times over time, showing fluctuations and potential anomalies.
- **Create Order Requests (Table):** A table with traces collected by **Jaeger**, with trace details shown on the side.

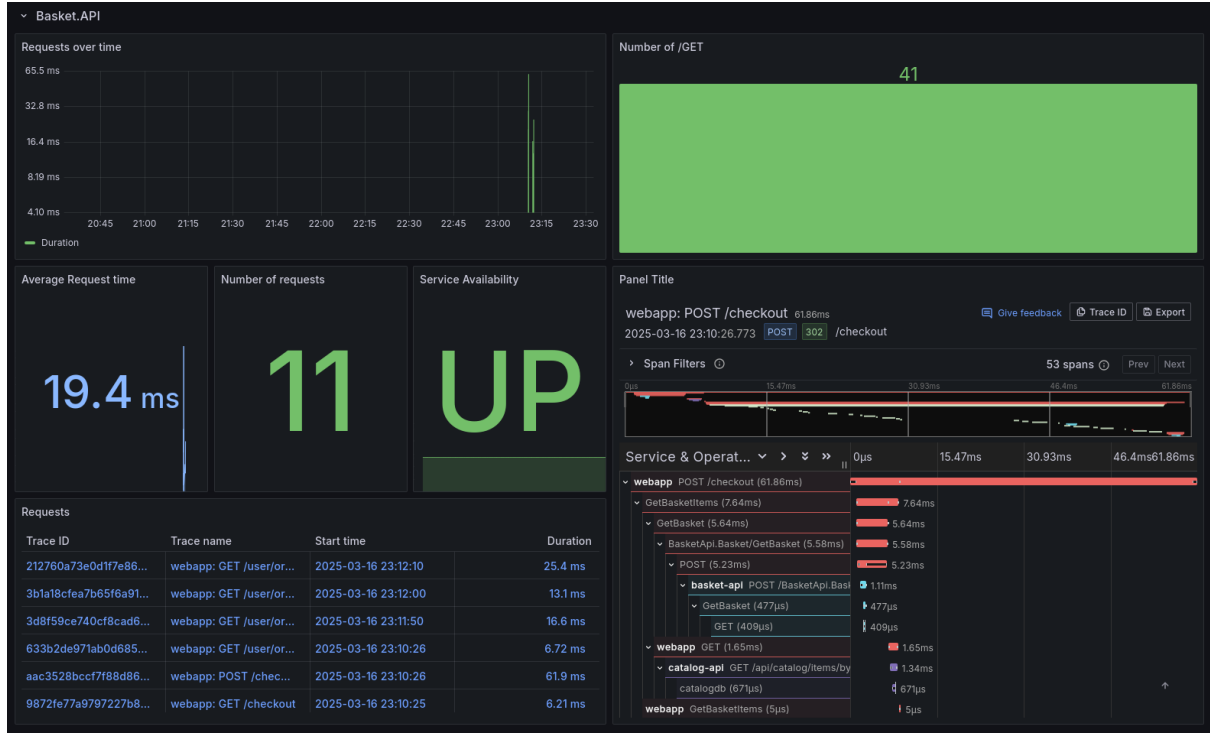


Figure 12: Basket Dashboard

Here, we display the following:

- **Requests over time:** A Time Series Graph that displays the API response times over time.
- **Number of /gets:** This metric comes from **Prometheus**, specifically from the meter that counts the number of /GET requests to the service.
- **Average request time:** The average time for each request to the Basket API.
- **Service availability:** This data comes from the health checks in **ASP.NET**.
- **Number of requests:** This metric counts the number of requests made to the Basket API, collected by **Jaeger**.
- **Requests (Table):** A table with traces collected by **Jaeger**, with trace details shown on the side.

3 Conclusion

In this report, I explored the process of implementing **observability** in a **microservices architecture** using tools such as **OpenTelemetry**, **Jaeger**, **Prometheus**, and **Grafana**. By instrumenting the microservices with **tracing** and **metrics collection**, we were able to gain valuable insights into the **performance** and **health** of the services, which are critical for maintaining a reliable and efficient application.

During the process, I encountered several challenges. The first challenge was understanding how the application was built and how it functioned. The second was learning

how to incorporate **OpenTelemetry**, **Prometheus**, and **Jaeger** with an app written in **.NET** and **C#**, tools I was not very familiar with. After overcoming this, I struggled with ensuring the communication between services and establishing proper connectivity with **Grafana**. One particular challenge was automating the saving of dashboards in **Grafana**, which, in the end, I was unable to achieve.

To help overcome these challenges, I leveraged **ChatGPT**, **Gemini**, and **Copilot**. These AI-powered tools assisted me in generating **configuration files**, exploring and navigating **C#** code, and troubleshooting issues I encountered during development. Additionally, **Copilot** helped me write code faster by suggesting relevant snippets and improving my coding efficiency.

In the end, I faced difficulties in creating **load tests** with **k6** to test the app due to time constraints, as well as explore data encryption and compliance in the **database** layer and introduce **column masking** for sensitive data, ending up not implementing these topics.

All the code developed, as well as the instructions to run it are available in my repository via this link (<https://github.com/Dan1m4D/AS-Assignment-1>).